

How to bring your Data Science Project in production

Using Azure Databricks with Spark, Azure ML and Azure DevOps



René Bremer

Follow

Jan 20, 2019 · 10 min read

1. Introduction

A lot of companies struggle to bring their data science projects into production. A common issue is that the closer the model is to production, the harder it is to answer the following question:

- Why did the model predict this?

Having a build/release pipeline for these data science project can help to answer this question. It enables you to trace back that:

- Model M was trained on dataset D with algorithm A by person P.
- Model M was deployed in production in release R on time T

This audit trail is essential for every model running in production and is required in a lot of industries, e.g. finance.

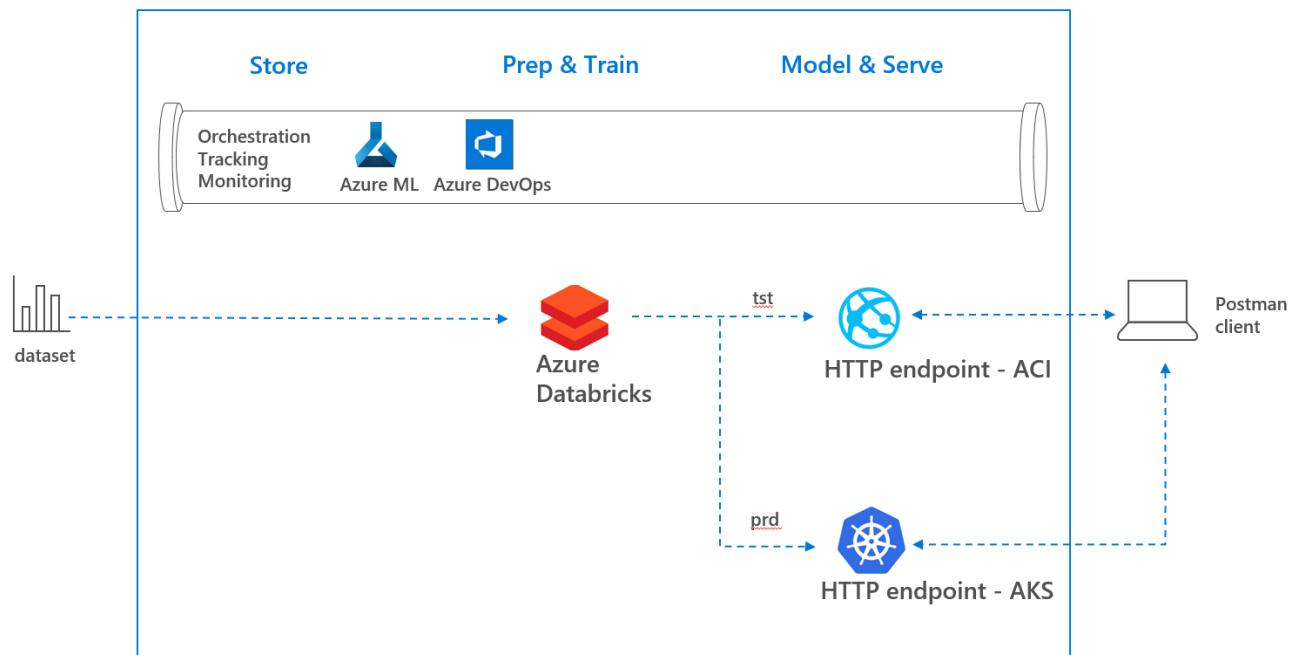
Last major update of blog/code on github: Jan 12, 2020

2. Objective

In this tutorial, a build/release pipeline for a machine learning project is created as follows:

- An HTTP endpoint is created that predicts if the income of a person is higher or lower than 50k per year using features as age, hours of week working, education.
- Azure Databricks with Spark, Azure ML and Azure DevOps are used to create a model and endpoint. Azure Container Instance (ACI) is used for testing, Azure Kubernetes Service (AKS) as production environment.

The project can be depicted in the following high level overview:



2. High level overview

In the remainder of this blog, the following steps will be executed:

- 3. Prerequisites
- 4. Create machine learning model in Azure Databricks
- 5. Manage model in Azure Machine Learning Service
- 6,7. Build and release model in Azure DevOps
- 8. Conclusion

The follow-up of the blog can be found here in which security is embedded in the build/release pipeline. Furthermore, the details of the audit trail are discussed in this

blog. Finally, if you interested how to use Azure Databricks with Azure Data Factory, refer to this blog.

3. Prerequisites

The following resources are required in this tutorial:

- Azure Databricks
- Azure Machine Learning Service
- Azure DevOps

4. Create machine learning model in Azure Databricks

Azure Databricks is an Apache Spark-based analytics platform optimized for Azure. It can be used for many analytical workloads, amongst others machine learning and deep learning. In this step, the following is done:

- 4a. Create new cluster
- 4b. Import notebook
- 4c. Run notebook

4a. Create new cluster

Start your Azure Databricks workspace and go to Cluster. Create a new cluster with the following settings:

Microsoft Azure

Create Cluster

New Cluster

Cancel

Create Cluster

Cluster Name

ml-test

Cluster Mode ?

Standard

Databricks Runtime Version ?

Runtime: 5.1 (Scala 2.11, Spark 2.4.0)

Learn more

1-2 Workers: 14.0-28.0 GB Memory, 4-8 Cores, 0.75-1.5 DBU
1 Driver: 14.0 GB Memory, 4 Cores, 0.75 DBU Cost \$0.55 per DBU ?

Python Version 3

Autopilot Options

☒ Enable autoscaling

☒ Terminate after 120 minutes of inactivity

Worker Type Standard_DS3_v2 14.0 GB Memory, 4 Cores, 0.75 DBU

Min Workers 1 Max Workers 2

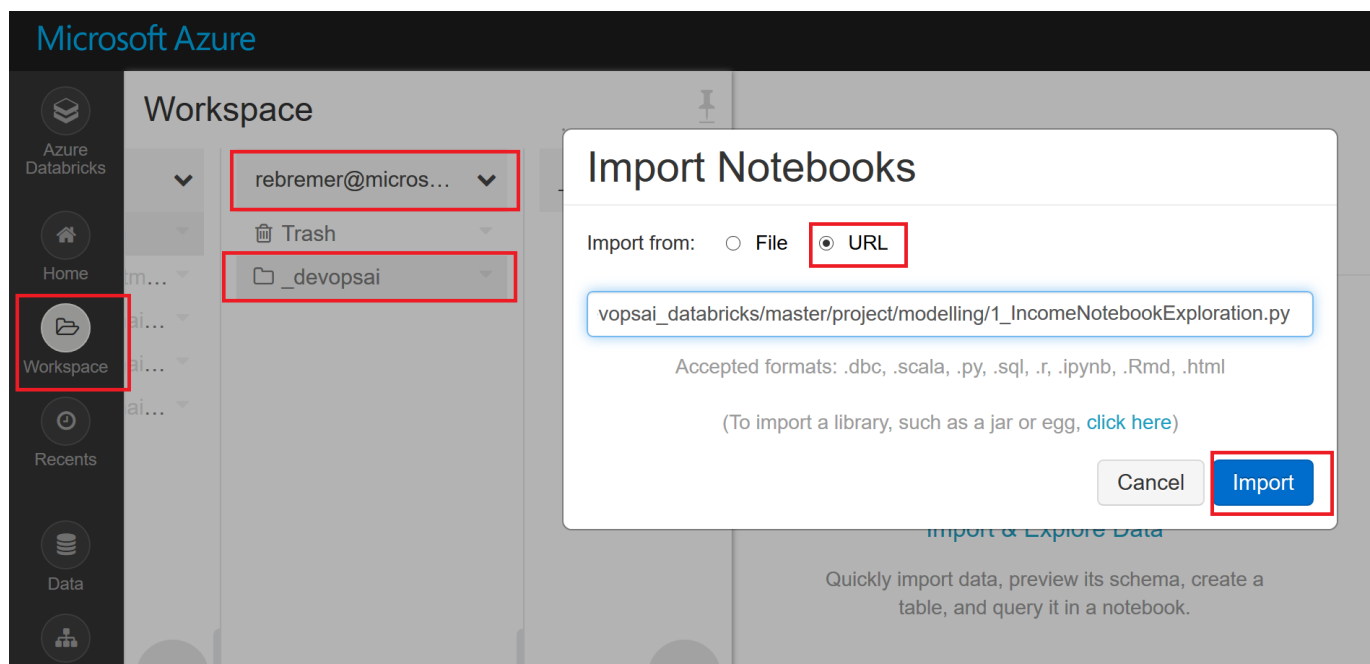
4a1. Create cluster

4b. Import notebook

Go to your Azure Databricks workspace, right-click and then select import. In the radio button, select to import the following notebook using URL:

https://raw.githubusercontent.com/rebremer/devopsai_databricks/master/project/modelling/1_IncomeNotebookExploration.py

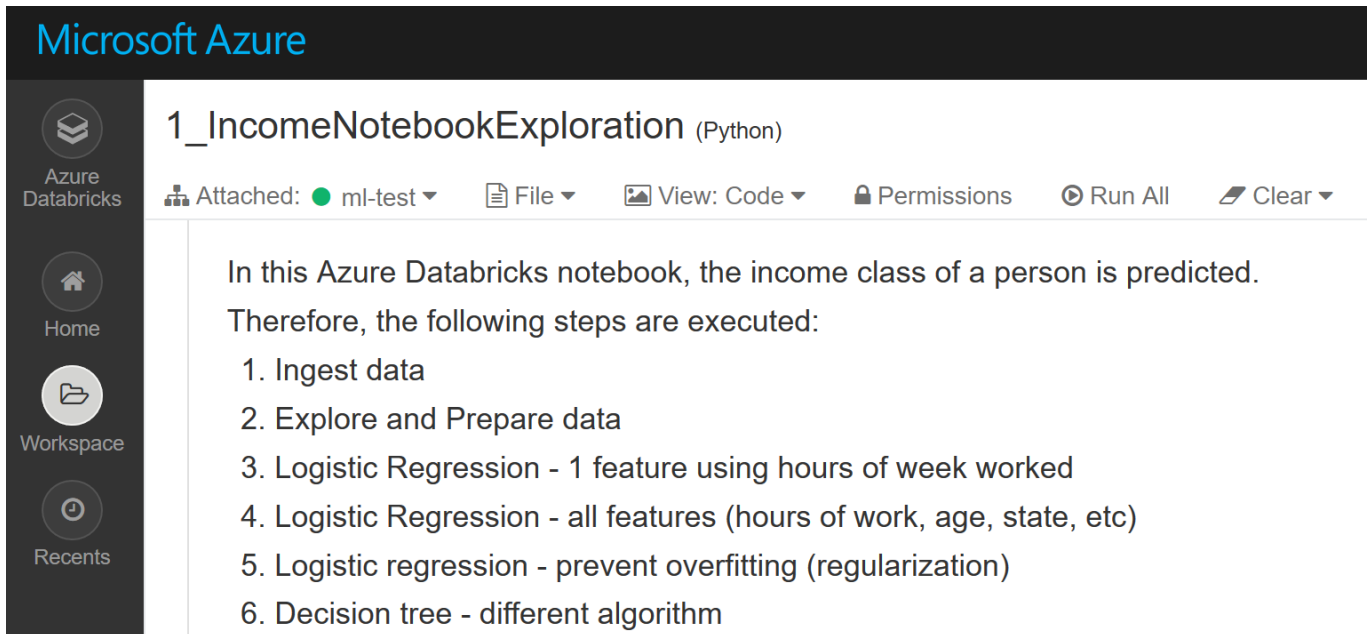
See also picture below:



4b1. Import notebook

4c. Run notebook

Select the notebook you imported in 4b and attach the notebook to the cluster you created in 4a. Make sure that the cluster is running and otherwise start it. Read the steps in the notebook, in which the data is explored and several settings and algorithms are tried to create a model that predicts the income class of a person. Walk through the notebook cell by cell by using shortcut SHIFT+ENTER.



4c1. Steps in notebook

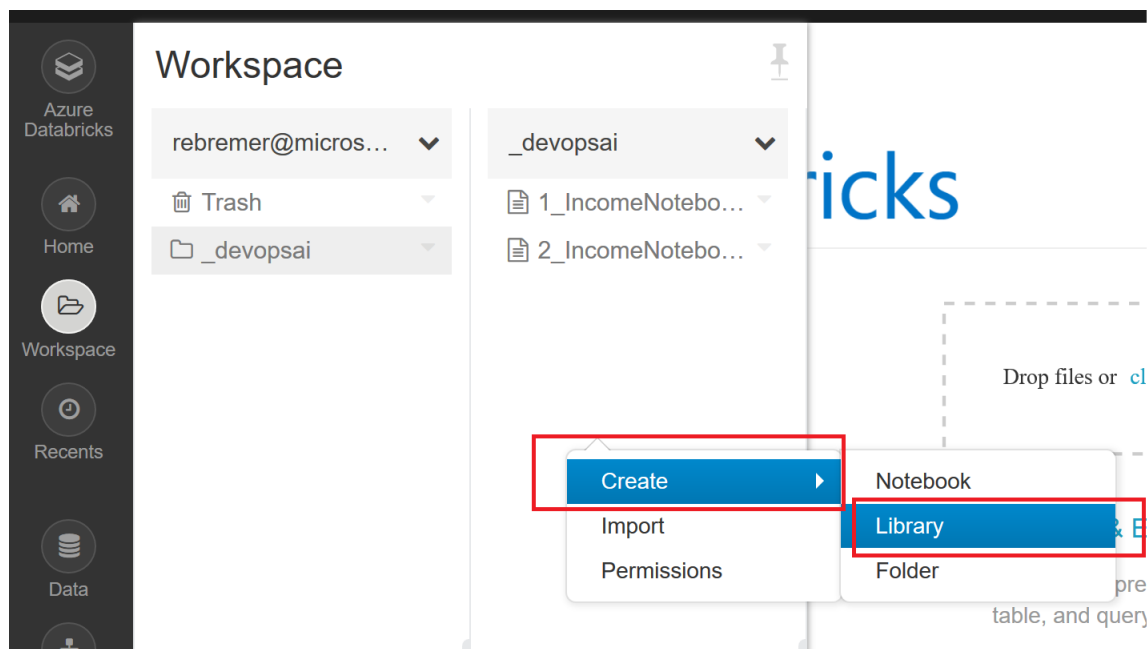
5. Manage model in Azure Machine Learning Service

Azure Machine Learning Service is a cloud service that you use to train, deploy, automate, and manage machine learning models. In this context, the model that was created in previous step will be added to your AMLS instance. The following steps will be executed

- 5a. Add library to Databricks cluster
- 5b. Import notebook using Azure ML to Azure Databricks
- 5c. Review results in Azure ML

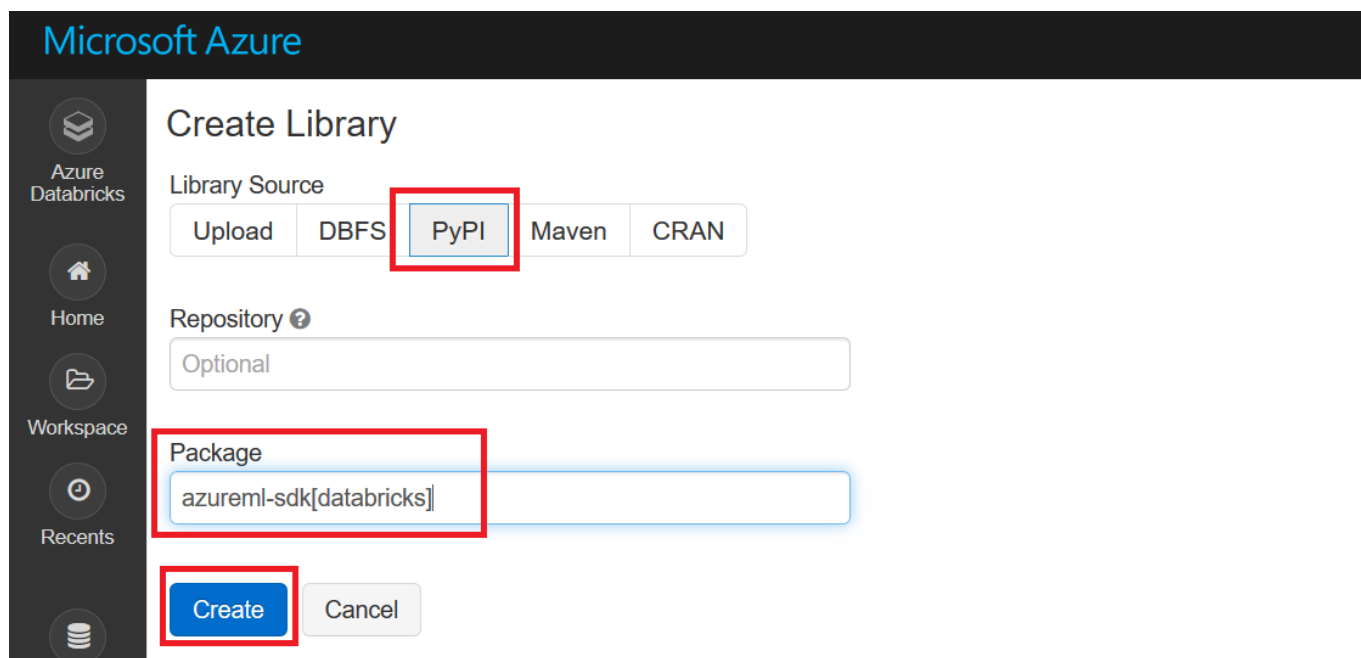
5a. Add library to Databricks cluster

Right click in your workspace and select to “create library”




5a1. Create library

Select PyPi and then fill in: azureml-sdk[databricks]



5a2. Add PyPi library with azureml-sdk[databricks]

Finally, attach the library to the cluster.



The screenshot shows the Azure Databricks interface. On the left is a sidebar with icons for Home, Workspace, Recents, and Data. The main area displays the 'azureml-sdk[databricks]' library page. It includes a 'Move to Trash' button, 'PyPI rules', and a link to 'azureml-sdk[databricks]'. Below this is the 'Status on running clusters' section, which has a checkbox for 'Install automatically on all clusters' (checked) and two buttons: 'Uninstall' and 'Install' (highlighted with a red box). A table below shows the installation status across clusters:

	Status	Cluster Name	Message
<input type="checkbox"/>	● Installed	ml-test	

5a3. Attach library to cluster

5b. Import notebook using Azure ML to Azure Databricks

In the previous part of this tutorial, a model was created in Azure Databricks. In this part you are going to add the created model to Azure Machine Learning Service.

Go to your Databricks Service again, right click, select import and import the a notebook using the following URL:

```
https://raw.githubusercontent.com/rebremer/devopsai_databricks/master/project/modelling/2_IncomeNotebookAMLS.py
```

Again, make sure it is attached to a cluster and the cluster is running

The screenshot shows the Azure Databricks workspace. The top bar displays 'Microsoft Azure'. The sidebar on the left has icons for Home, Workspace (highlighted with a red box), Recents, and Data. The main area shows the notebook '2_IncomeNotebookAMLS (Python)'. Above the notebook content, there is a toolbar with 'Attached: ● ml-test' (highlighted with a red box), 'File', 'View: Code', 'Permissions', 'Run All', and 'Clear'. The notebook content includes the following text:

```
Cmd 1

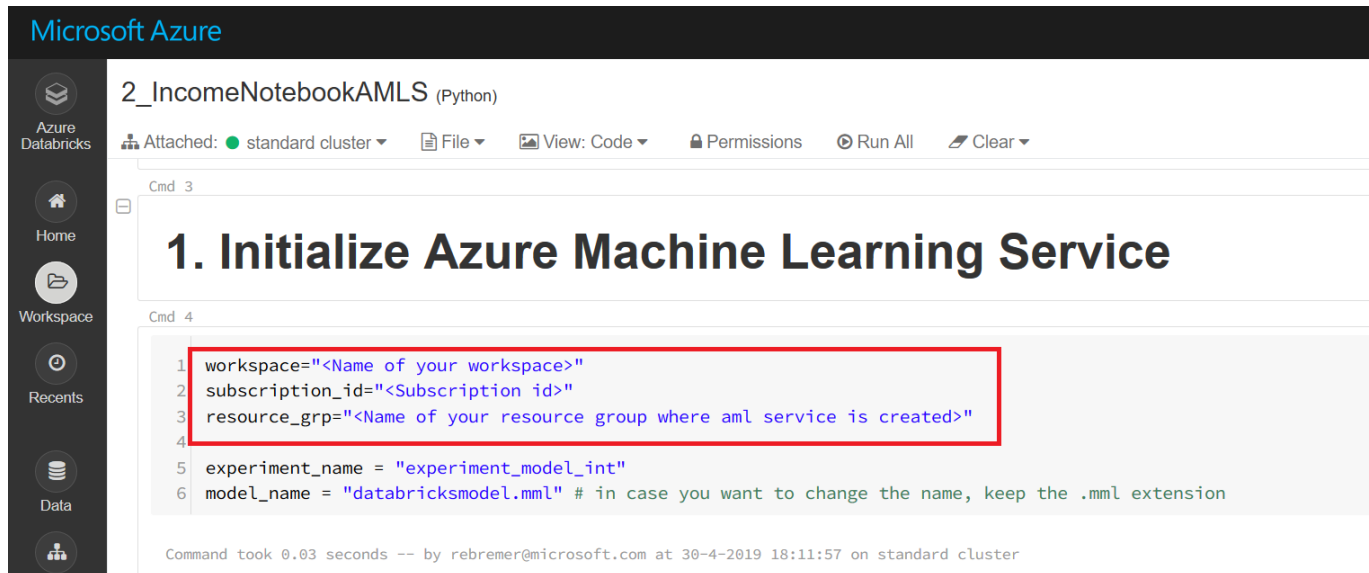
Azure ML & Azure Databricks notebooks by Rene Bremer (original taken from Parashar Shah)

Copyright (c) Microsoft Corporation. All rights reserved.

Licensed under the MIT License.
```

5b1. Import AMLS notebook

Subsequently, fill in the correct values for workspace, subscription_id and resource_grp. All values can be found in the overview tab of your Azure Machine Learning Service Workspace in the Azure Portal.



5b2. Add variables notebook

Now run the notebook cell for cell by using shortcut SHIFT+ENTER.

In cell 6, you will need to authenticate to Azure Machine Learning Service in the notebook. Follow the instruction in the notebook by opening the URL and enter the generated code to authenticate.

5c. Review results in Azure ML

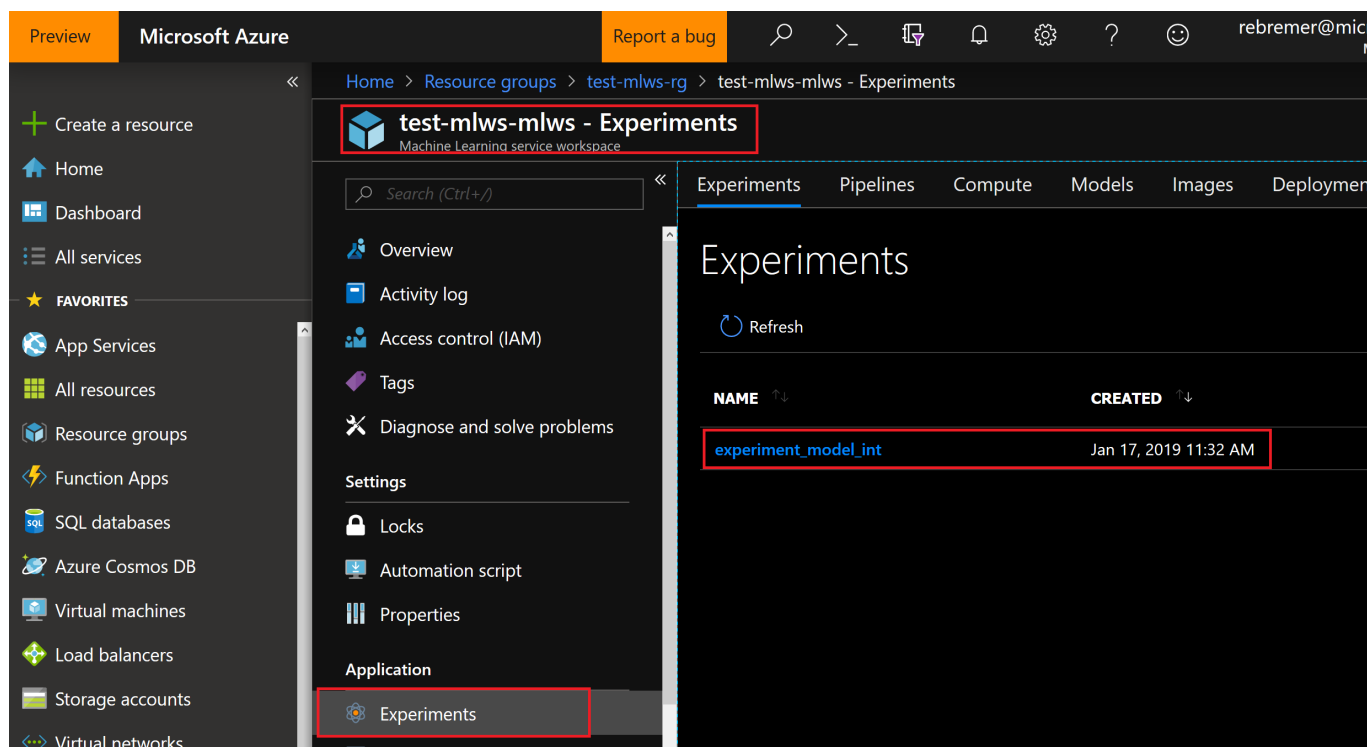
In step 5b, a notebook was run in which the results were written to Azure Machine Learning Service. In this, the following was done:

- A new experiment was created in you Azure ML
- With in this experiment, a root run with 6 child runs were the different attempts can be found.
- A childrun contains a description of the model (e.g. Logistic Regression with regularization 0) and the most important logging of the attempt (e.g. accuracy,

number of false positives)

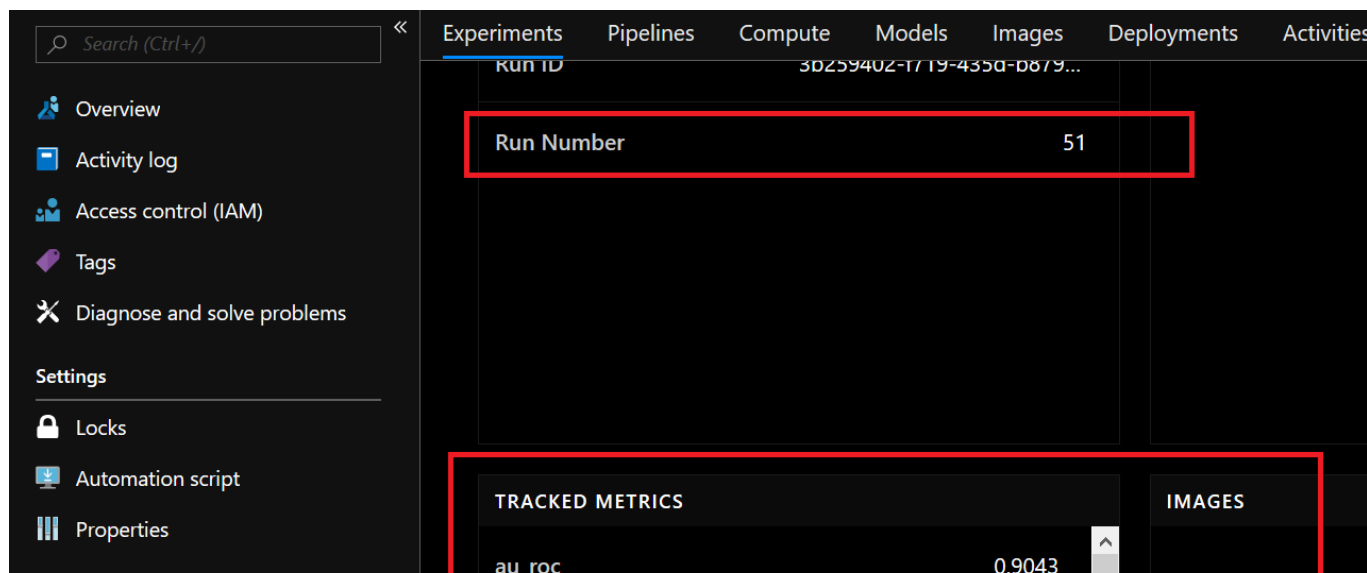
- The model artifact (.mml) is also part of a childrun. The artifact of the best childrun can be taken and deployed into production.

Go to you Azure ML instance. Select the experiment name that was used in the notebook (e.g. experiment_model_int).



5c1. Find experiment in Azure Machine Learning Service

Now click on the experiment, click on the run and childrun you want to see the metrics.



Application	
Experiments	
Pipelines	
Compute	
Models	
Images	
Deployments	
Activities	

au_prc	0.7611
truePositive	1416
falsePositive	500
trueNegative	6895
falseNegative	942

5c2. Find metrics of a childrun in Azure Machine Learning Service

When you go to output, you will find the model artifact, which you can also download. The model artifact of the best run will be used as the base of the container that is deployed using Azure DevOps in the next part of this tutorial.

test-mlws-mlws - Experiments
Machine Learning service workspace

Search (Ctrl+/)

Overview
Activity log
Access control (IAM)
Tags
Diagnose and solve problems

Settings
Locks
Automation script
Properties

Application
Experiments

Experiments Pipelines Compute Models Images Deployments Activities

experiment_model_int | Run Number: 51

Back to Parent Run Refresh Get Link

Details **Outputs** Logs Snapshot

Outputs

NAME	DOWNLOAD
databricksmodel.mml	Download

5c3. Model artifact

6. Create and prepare Azure DevOps project

Azure DevOps is the tool to continuously build, test, and deploy your code to any platform and cloud. In chapter 6, an Azure DevOps will be created and prepared. The project will be prepared using the following steps:

- 6a. Create Personal Access Token in Databricks
- 6b. Create AKS cluster (optional)

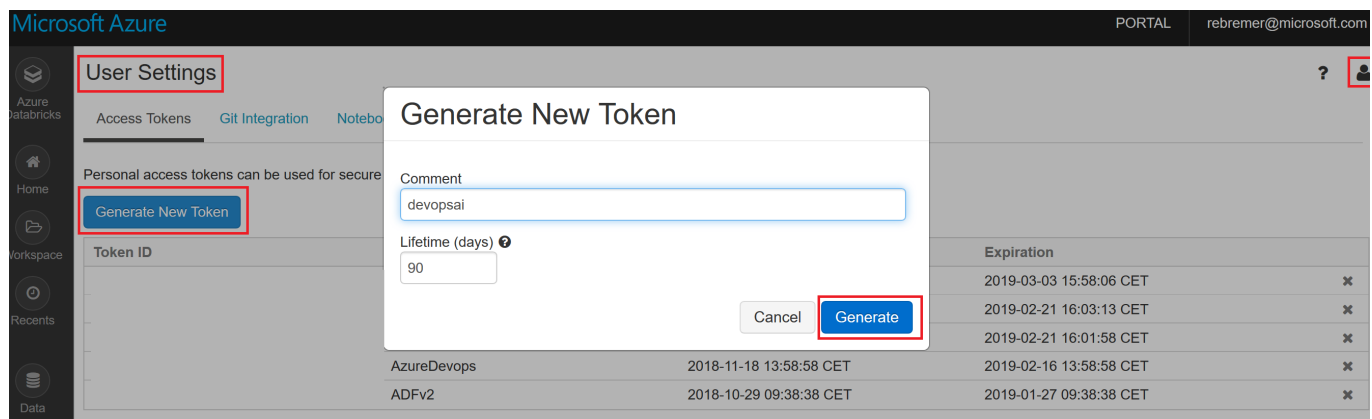
- 6c. Create Azure DevOps project and service connection
- 6d. Add variables to code

In chapter 7, the actual build-release pipeline will be created and run to create an endpoint of the model.

6a. Create Personal Access Token in Databricks

To run Notebooks in Azure Databricks triggered from Azure DevOps (using REST APIs), a Databricks Access Token (PAT) is required for authentication.

Go to Azure Databricks and click to the person icon in the upper right corner. Select User Settings and then generate a new token.




6a1. Generate Databricks Access Token

Make sure to copy the token now. You won't be able to see it again. Token is needed to access Databricks from the Azure DevOps build pipeline later

6b. Create AKS cluster (optional)

Typically, ACI is used for testing, whereas webapps or AKS cluster is used for a production situation. In this optional step, an prd environment is created in which the container is deployed to an AKS cluster. First, go to to you Azure ML Service Workspace and select Compute. Take as compute name **blog-devai-aks** and select Kubernetes Service as compute type, see also below.



Add Compute

- * Compute name:
- * Compute type:
- * Kubernetes Service: ☒ Create new ☐ Use existing
- * Region:
- * Virtual machine size:
[Choose Virtual machine size](#)
- * Number of nodes:
- * Network configuration: ☒ Basic ☐ Advanced

6b1. Add AKS compute to Azure ML Service

Creating an AKS cluster takes approximately 10 minutes. Continue to the next step.

6c. Create Azure DevOps project with service connection

Create a new project in Azure DevOps by following this tutorial. Once you create a new project, click on the repository folder and select to import the following repository:

https://github.com/rebremer/devopsai_databricks.git

See also picture below:

rebremer / devopsai databricks final / Repos / Files / devopsai databricks final

devopsai databricks final is empty. Add some code!

Clone to your computer

HTTPS SSH OR

Having problems authenticating in Git? Be sure to get the latest version of [Git for Windows](#) or our plugins for [IntelliJ](#), [Eclipse](#), [Android Studio](#) or [Windows command line](#).

or push an existing repository from command line


HTTPS SSH

git remote add origin https://rebremer.visualstudio.com/devopsai%20databricks%20final/_git/...

git push -u origin --all

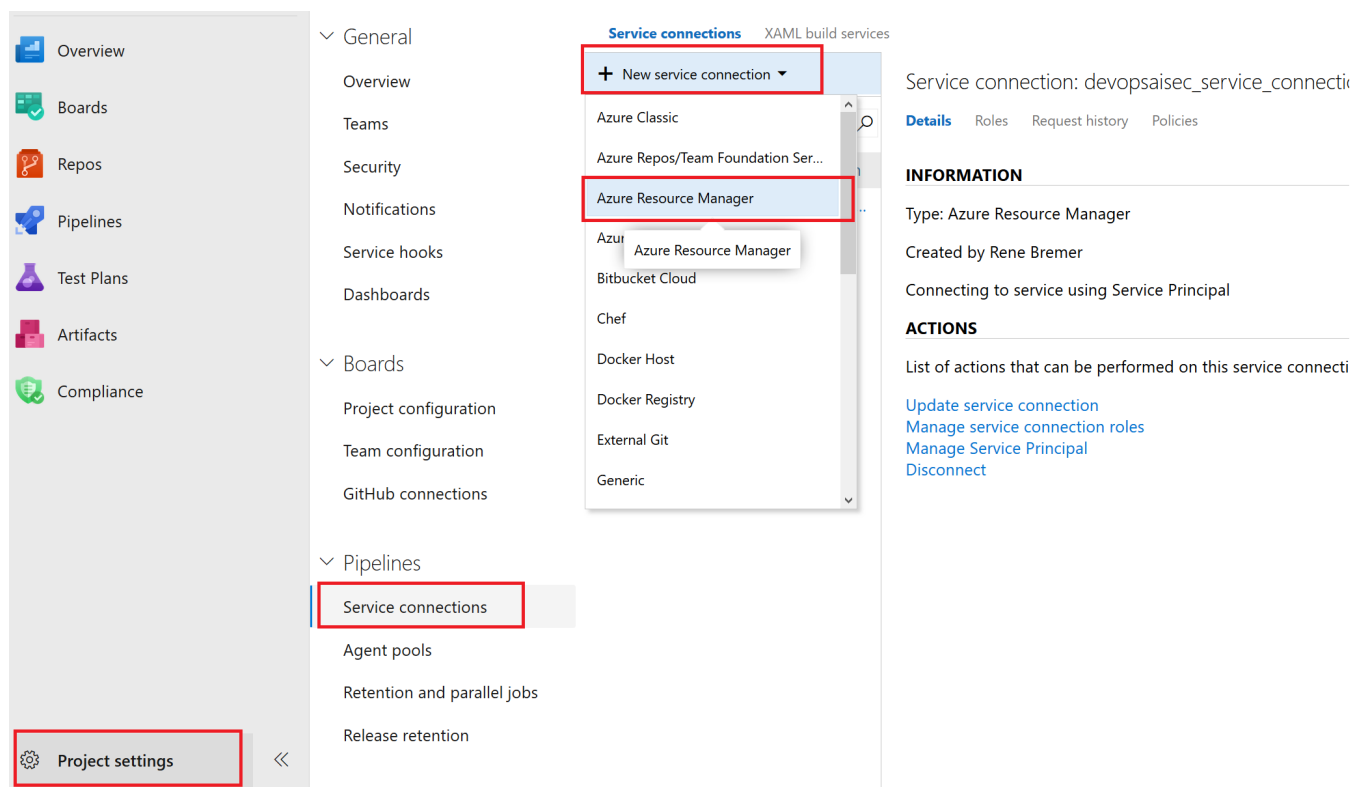
Import a Git repository

Source type



6c1. Add repository to your Azure DevOps project

A Service connection is needed to access the resources in the resource group from Azure DevOps. Go to project settings, service connection and then select Azure Resource Manager.



6c2. Go to Service Connection

Select Service Principal Authentication and limit scope to your resource group in which your Machine Learning Workspace Service is deployed. Make sure that you name the connection as follows: **devopsaisec_service_connection**.

Add an Azure Resource Manager service connection

☒ Service Principal Authentication

☐ Managed Identity Authentication

Connection name	<input type="text" value="devopsai-serviceconnection"/>
Scope level	<div>Subscription</div>
Subscription	<div>Microsoft Azure Internal Consumption (513a7987-bc</div>
Resource Group	<div>devaisec-rg</div>

Subscriptions listed are from Azure Cloud

6c3. Create Azure Resource Manager service connection

6d. Add variables to code

In the Repos you created in the previous step, the following files shall be changed:

- \project\configcode_build_release.yml

With the same variables for workspace, subscription_id and resource with values of your Machine Learning Service Workspace as in step 5b. Also, fill in your Databricks Personal Access Token generated in step 6a.

```
variables:  
  # change 5 variables below with your own settings, make sure that  
  # : with a space is kept and not replaced with =  
  workspace: '<<Name of your workspace>>'  
  subscription_id: '<<Subscription id>>'  
  resource_grp: '<<Name of your resource group with aml service>>'  
  domain: 'westeurope.azuredatabricks.net' # change loc.when needed  
  dbr_pat_token_raw: '<<your Databricks Personal Access Token>>'
```

The files can be changed in Azure DevOps by looking up the file in the Repos, click on “edit”, change the variables and then “commit” the file. You can also clone the project and work from there. Notice that in a production situation, keys must never be added to a code. Instead, secret variables in an Azure DevOps pipeline shall be used and is dealt with in this follow-up tutorial.

In this chapter, an Azure DevOps project is created and prepared. Now the model is ready to be built and released in the Azure DevOps project.

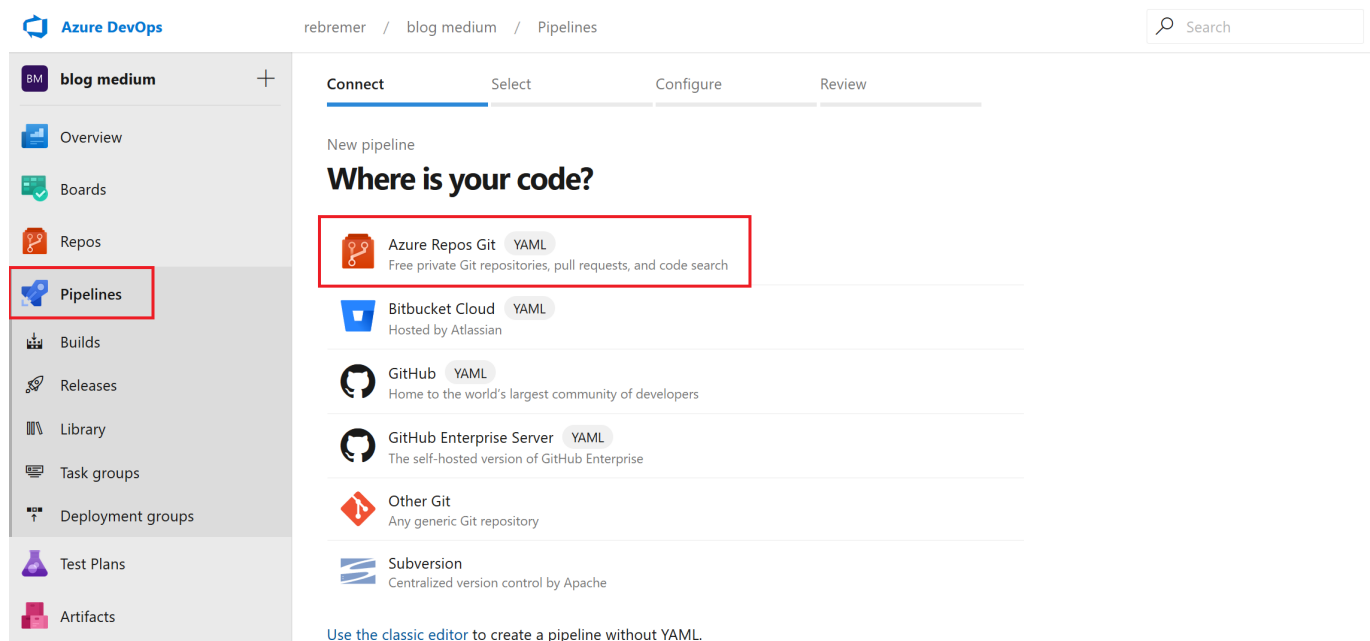
7. Build and release model in Azure DevOps

In this part, the model is built and released in the Azure DevOps using the following steps:

- 7a. Create build-release pipeline
- 7b. Run build-release pipeline
- 7c. Consume HTTP endpoint with Postman

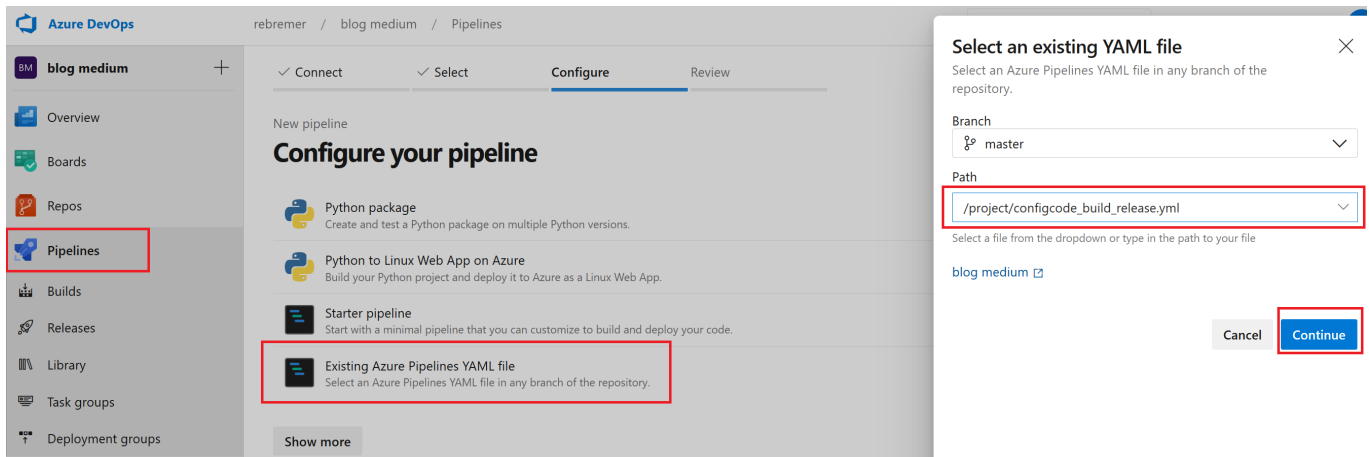
7a. Create build-release pipeline

In this step, you are going to create a build-release pipeline. Go to Azure DevOps project you have created in 6c and then click on Pipelines. A wizard is shown in which your Azure Repos Git shall be selected, see also below.



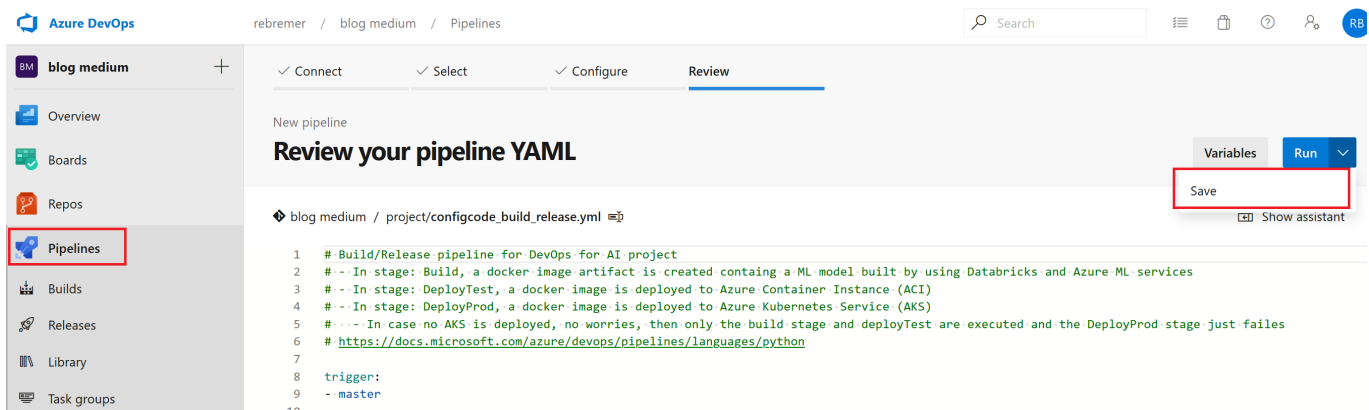
7a1. Create Pipeline

Subsequently, select your Git repo attached to this project and then select “Existing Azure Pipelines YAML file”. Then browse the directory “project/configcode_build_release.yml”, see also below.



7a2. Select build-release YAML file

Finally review your pipeline and save your pipeline, see also below.



7a3. Save pipeline

In this chapter, the pipeline was configured. In this pipeline the following steps will be executed:

Build:

- Select Python 3.6 and install dependencies
- Upload notebook to Databricks
- Create model using Azure Databricks by running notebook. Add model to Azure Machine Learning service
- Create a docker image from model that will be deployed in release step
- Creation of build artifact as input for release deployTest and deployProd

Release deployTest:

- Retrieve docker image created in Build step
- Deploy docker image to Azure Container Instance (ACI)
- Test endpoint in ACI

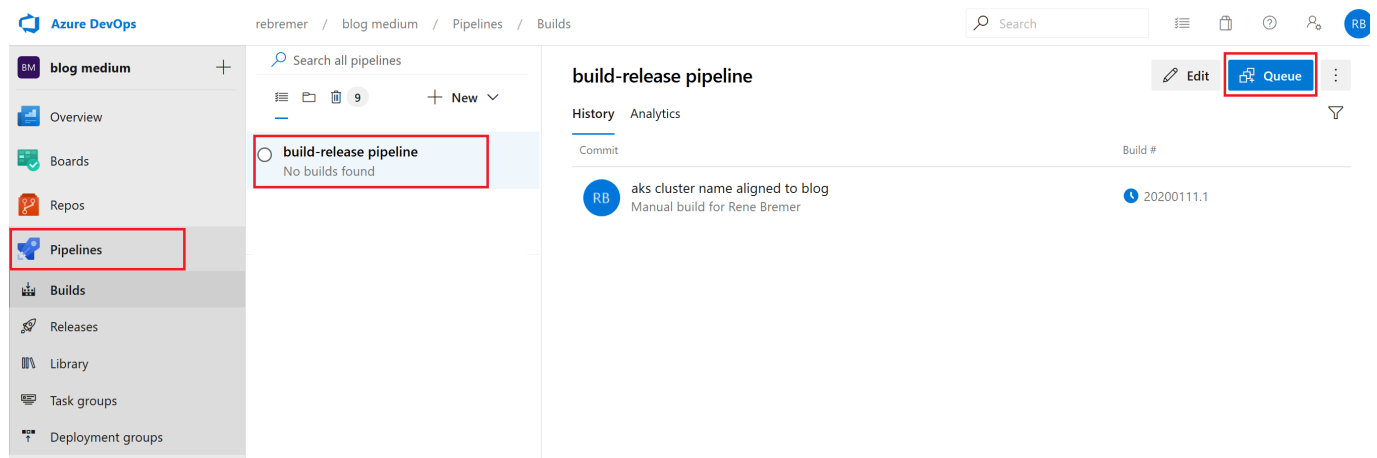
Release deployProd (optional, when AKS cluster is deployed in step 7a):

- Retrieve docker image created in Build step
- Deploy docker image to Azure Kubernetes Service (AKS)

In the next part, the pipeline will be run.

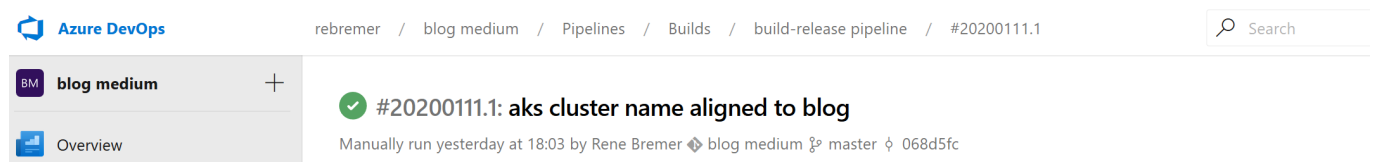
7b. Run build-release pipeline

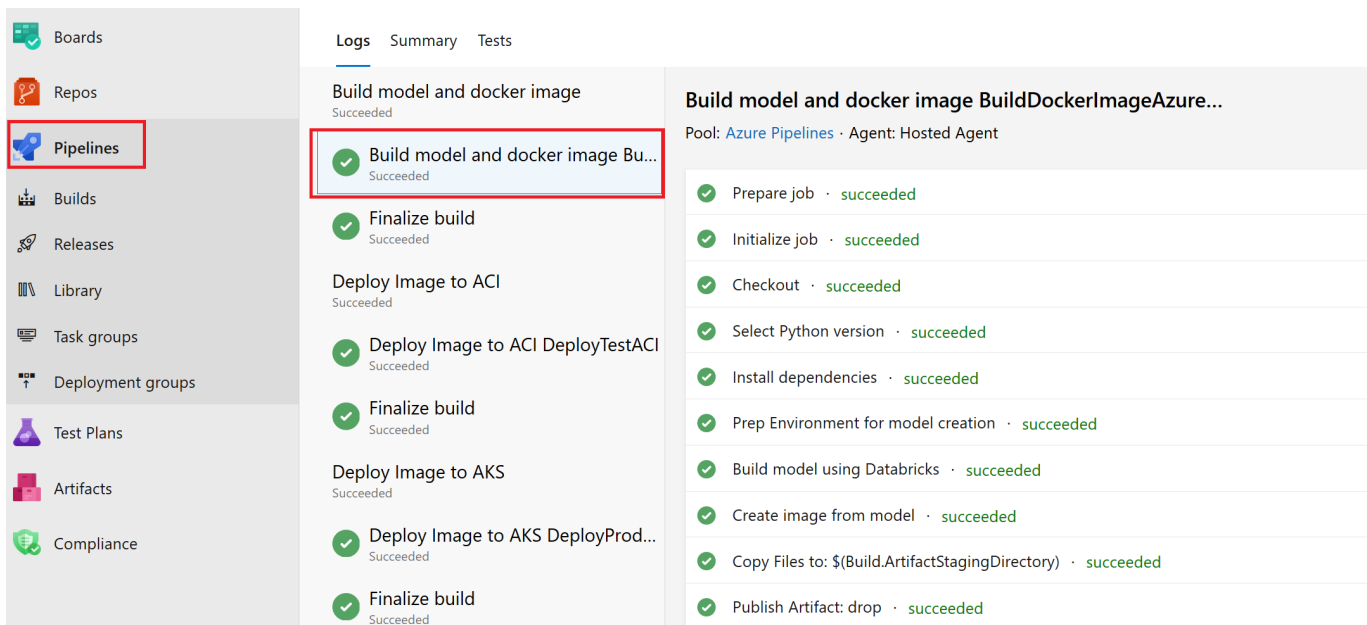
In this step, the build-release pipeline will be run in Azure DevOps. Go to your pipeline deployed in the previous step, select the pipeline and then select queue, see also below.



7b1. Run build-release pipeline

When the pipeline is started, a docker image is created containing an ML model using Azure Databricks and Azure ML in the build step. Subsequently, the docker image is deployed/released in ACI and AKS. A successful run can be seen below.



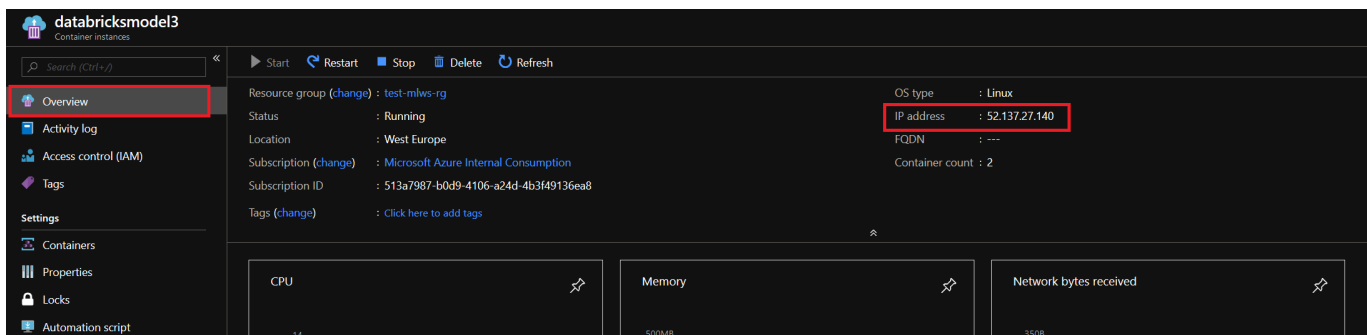


7b2. Successful run of build-release pipeline

Notice that if you decided to not deploy the docker image in AKS, the previous steps will still be executed and the AKS step will fail. For detailed logging, you can click on the various steps.

7c. Consume HTTP endpoint with Postman

When you go to the Azure Portal, you can find the Azure Container Instance deployed in 7b and subsequently, you can find the public IP address of the HTTP endpoint.

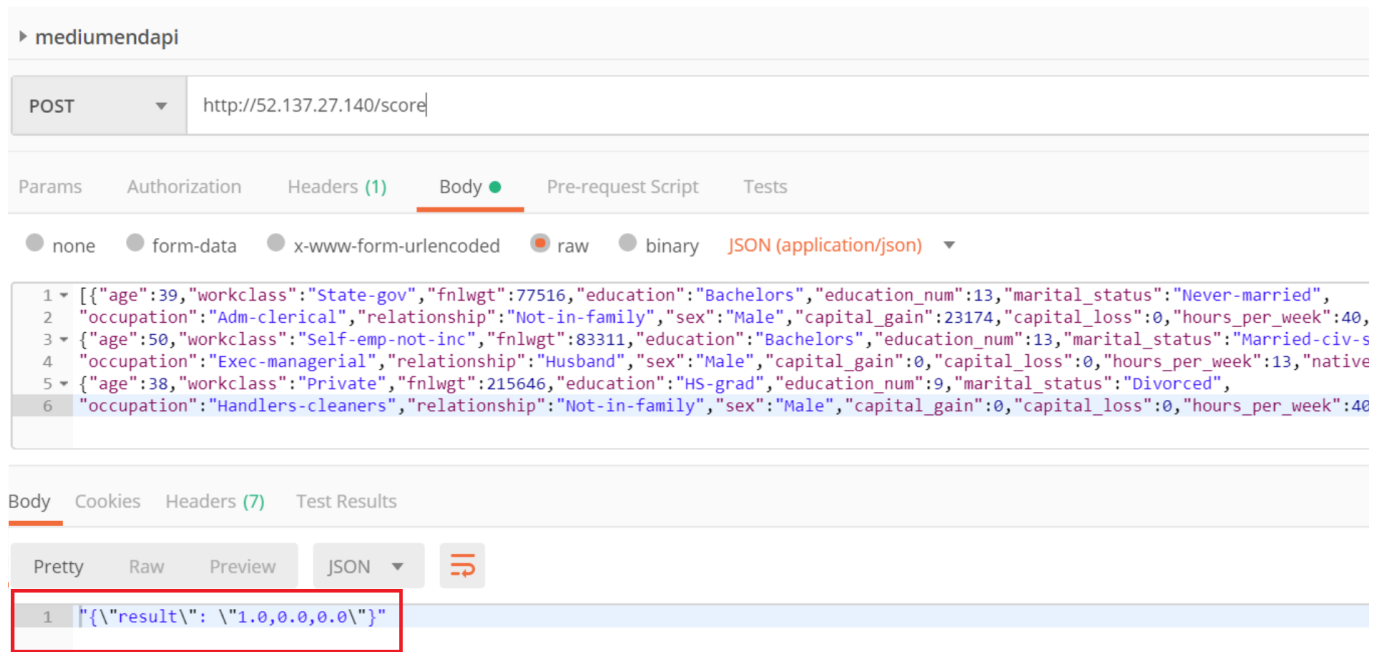


7c1. IP address of HTTP endpoint

This IP address will now be consumed by Postman to create predictions. An example payload can be found in the project/services/50_testEndpoint.py in the project. In this example, the income class of three persons is predicted.

- The prediction for the first person is that the income is higher than 50k,

- For the other two persons the prediction is lower than 50k.



7c2. Use HTTP endpoint to create predictions

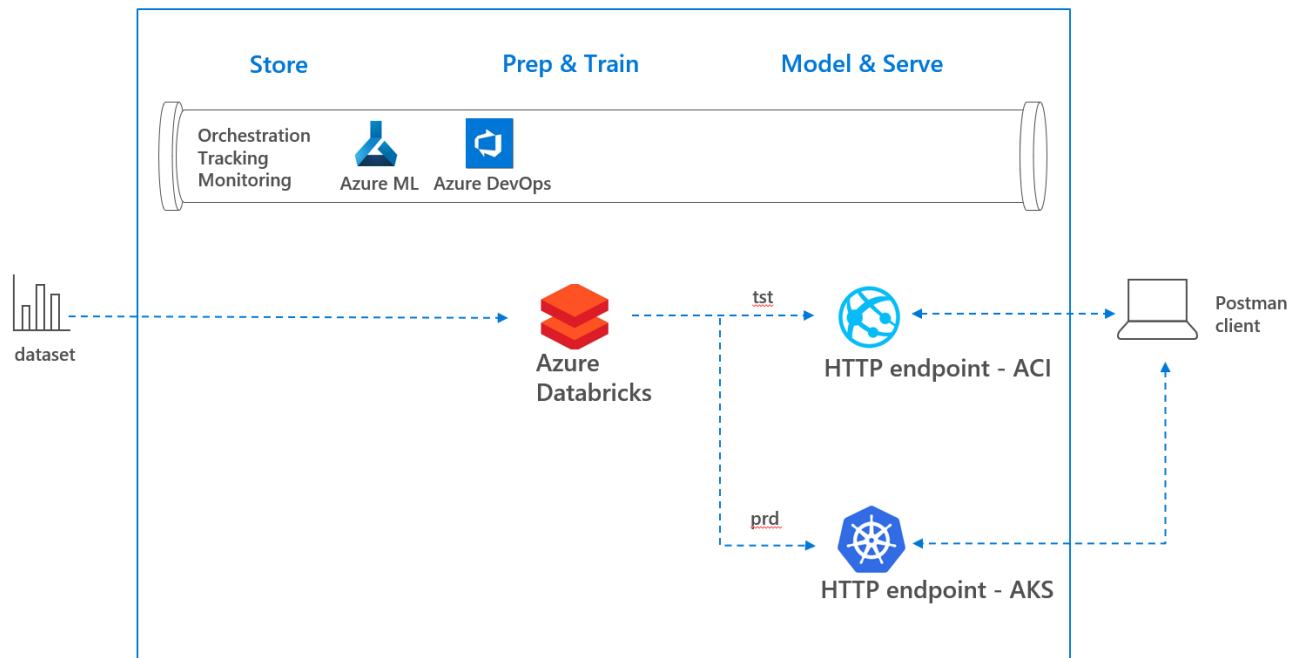
8. Conclusion

In this tutorial, an end to end pipeline for a machine learning project was created. In this:

- Azure Databricks with Spark was used to explore the data and create the machine learning models.
- Azure Machine Learning Service was used to keep track of the models and its metrics.
- Azure Devops was used to build an image of the best model and to release it as an endpoint.

This way you can orchestrate and monitor the entire pipeline from idea to the moment that the model is brought into production. This enables you to answer to question: *Why did the model predict this?*

The architecture overview can be found below. In this follow-up tutorial, security of the pipeline is enhanced.



8. High level overview

[Azure](#) [Azure Databricks](#) [Continuous Integration](#) [Data Science](#) [Machine Learning](#)

[About](#) [Help](#) [Legal](#)